

DAT103

Datamaskiner og operativsystemer

Processes

Chapter 3 in [B1]

Violet.Ka.I.Pun@hvl.no



Høgskulen
på Vestlandet

What do we cover?

- ▶ What a process is
- ▶ Characteristics of processes
 - Scheduling
 - Creation
 - Termination
 - Communication
- ▶ Communication with shared memory and messages
- ▶ Communication in client-server systems

Process concept (1)

- ▶ OS runs many different programs
 - Batch jobs
 - Time shared programs or tasks
- ▶ The terms **jobs** and **processes** are used interchangeably


Program is a passive entity on the disk, while **process** is active

- ▶ A program becomes a process when it is loaded into the memory
- ▶ A **process** is a program in execution
- ▶ Process execution is **sequential**

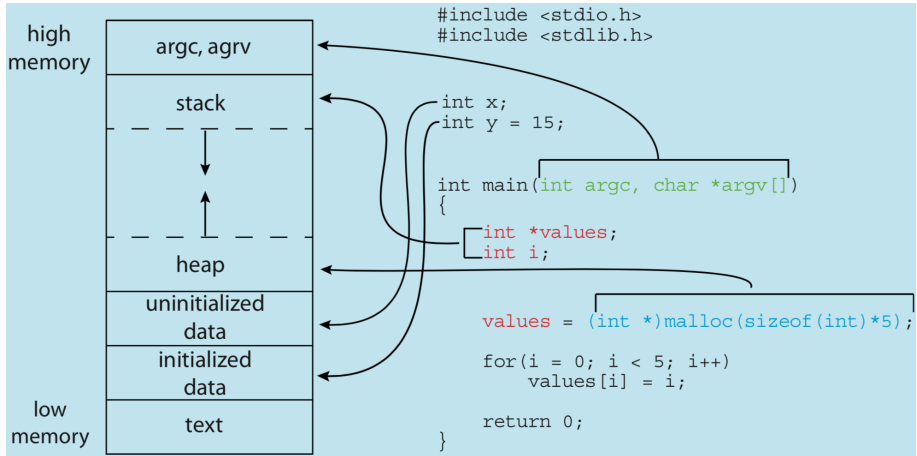
Process concept (2)

A process contains several parts

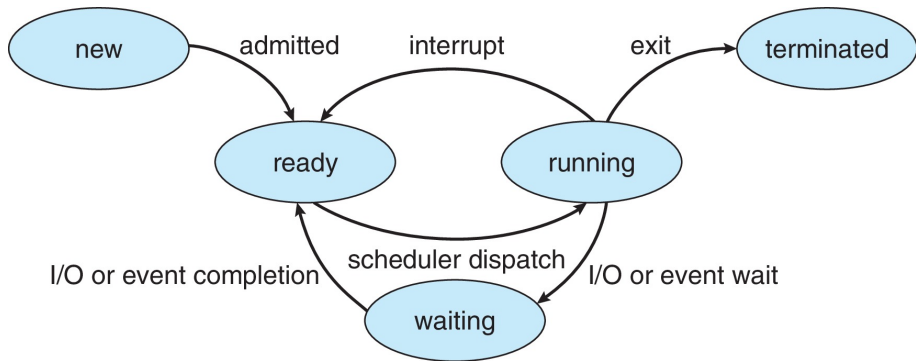
- ▶ **Program code:** called **text section**
- ▶ **Current activity:** **program counter** and content of processor's **registers**
- ▶ **Process stack:** temporary data (parameters, return address, local variables)
- ▶ **Data section:** global variables
- ▶ **Heap:** dynamically located **memory** during runtime
- ▶ Execution starts via mouse-click and command line
- ▶ A process itself can be an execution environment for other code



Example: Memory of a process for a C program




- ▶ A process can change its **state** during its execution
- ▶ The state of a process is defined partly by the current activity of that process
- ▶ A process may be in one of these states
 - **New**
 - **Running**
 - **Waiting**
 - **Ready**
 - **Terminated**
- ▶ Only **one** process can be **running** on any processor at anytime



Process control block (PCB)

Or, task control block

- ▶ Represents a process in OS
- ▶ Contains information associated with the process
- ▶ Includes
 - Process state
 - Program counter
 - CPU registers
 - CPU-scheduling information
 - Memory-management information
 - Accounting information
 - I/O status information

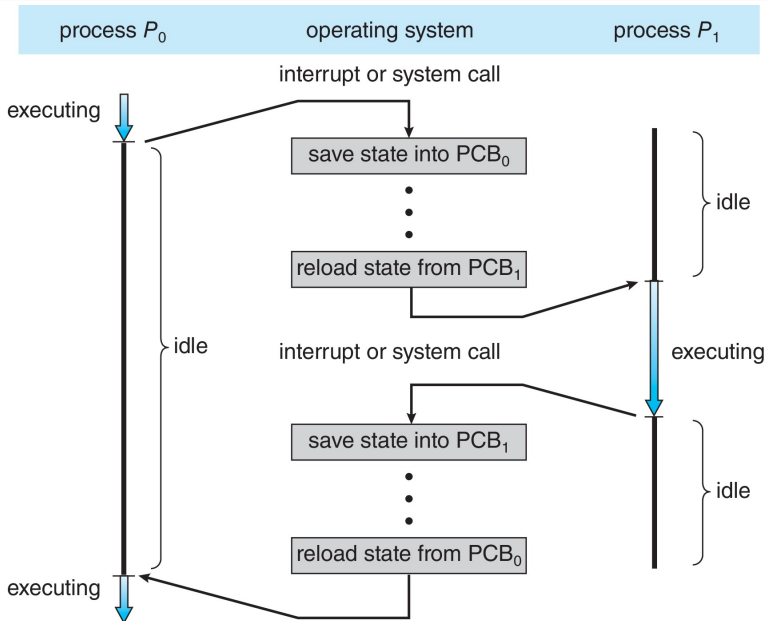


- ▶ Each process in the previous slides are assumed to have **only one thread**
- ▶ With **multiple threads** in a process
 - Each thread has a different program counter
 - Each program counter must be saved in PCB
- ▶ Must also include thread information in PCB

Thread Control Block (TCB)

- ▶ Contains information about:
 - Thread Identifier
 - CPU registers, including stack pointer, program counter, etc.
 - Thread state
 - Pointer to the Process control block (PCB) of the process to which the thread belongs

Switching CPU from process to process

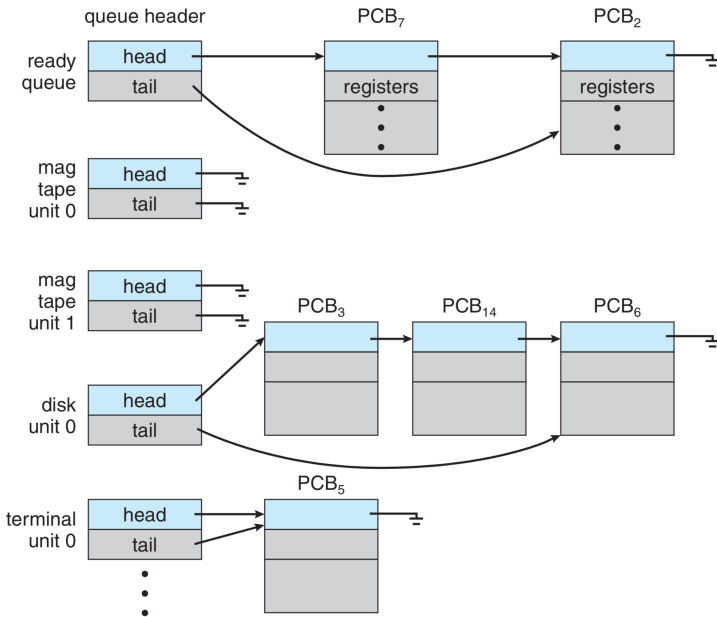


Context switch

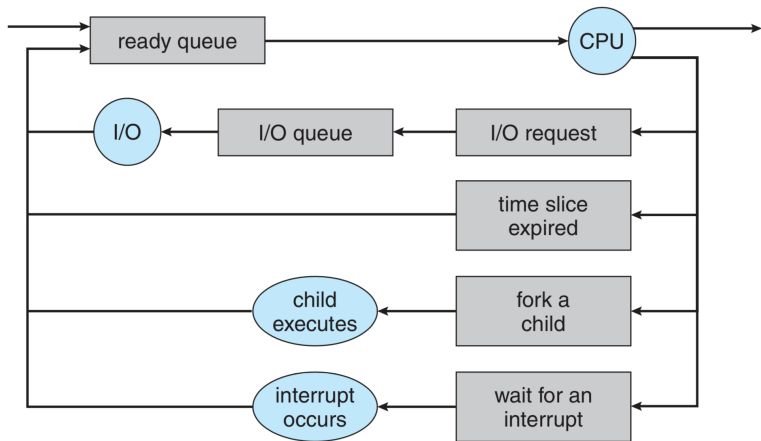
- ▶ When CPU switches process, the system must save the **state** of the current process and load the state of the new one
 - Need **context switch**
- ▶ Context is represented in PCB
- ▶ Time for context switch creates **overhead**
 - The system does not do useful work while switching
 - Switching speed varies from machine to machine
 - Factors may include: the memory speed, the number of registers that must be copied, ...
 - Typical speed: a few milliseconds instructions
 - The more complex the OS and the PCB
 - the longer the context switch
- ▶ Time taken also depends on hardware
 - Some hardware allow loading multiple contexts at once

- ▶ Maximise CPU utilisation by switching CPU among processes
- ▶ **Process scheduler** selects one out of all available processes for execution on CPU
- ▶ All other processes have to stay in some **queues**
 - **job queue**: all processes in the system
 - **ready queue**: processes in main memory, ready and wait to execute
 - usually stored as a linked list
 - **device queue**: processes waiting for I/O device
 - each device has its own device queue

Example: ready queue and device queues



Queueing-diagram representation of process scheduling



circles: resources

Schedulers (1)

- ▶ To select processes from various queues appropriately
- ▶ **Long-term scheduler** (job scheduler)
 - Select processes for **disk** to loads into **memory**
 - Control the **degree of multiprogramming**
i.e., the number of processes in the main memory
 - Selection is perform **less frequently** and can be **slow**
- ▶ **Short-term scheduler** (CPU scheduler)
 - Select a process from the **ready queue** and allocate the **CPU**
 - Selection is perform **frequently** and needs to be **fast**

Long-term scheduler


- ▶ Needs to select a good mix of I/O-bound and CPU-bound processes

Process can be described as either

- ▶ I/O bound, or
 - More time for I/O and computations
- ▶ CPU bound
 - Mostly computations, infrequent I/O requests

Mid-term scheduler

- ▶ Can remove a process from memory/CPU
→ reduce the degree of multiprogramming
- ▶ Can reintroduce the process into memory
- ▶ Called **swapping**¹
- ▶ Are used to improve the process mix; or
- ▶ To free up memory due to e.g., a change in memory requirements that has overcommitted available memory



¹Discussed later in Chapter 8

- ▶ Some mobile systems (e.g., early version of iOS) allow only one process to run at one time, while others suspended
- ▶ **iOS** currently allows:
 - **Single foreground** process: the one that is open and on the display
 - **Multiple background** processes: runs in memory and not on the display, and with limits
 - Background processes are limited to
 - single, short task, receiving notification of events and specific long-running tasks like audio player
- ▶ **Android** has fewer limits on background processes

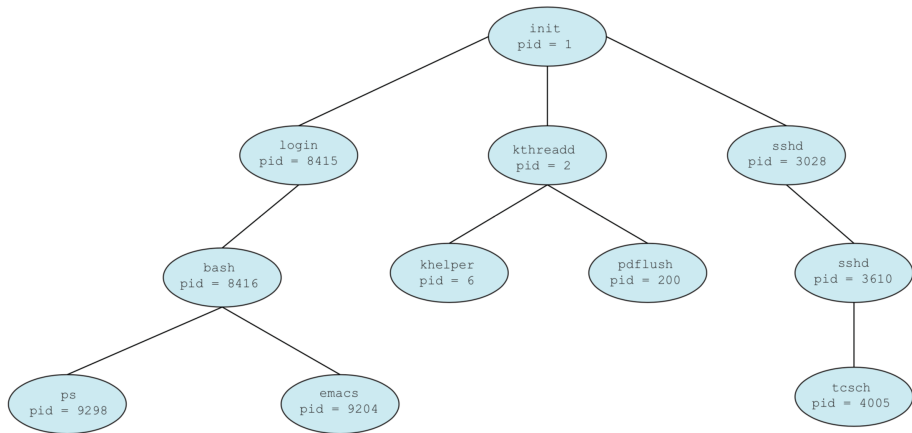
Operations on processes – Creation

- ▶ A **parent** process creates **child** processes
- ▶ Identify and manage a process with a **unique process identifier** (pid)

Relationship between parent and child processes

- ▶ Resource sharing
 - Parent and children share all resources, or
 - Children share subset of parent's resources, or
 - Parent and children do not share resources
- ▶ Execution
 - Parent and children run simultaneously, or
 - Parent waits until children are finished
- ▶ Address space
 - Child process is a duplicate of the parent process, or
 - Child process has a new program loaded into it

A typical process tree for Linux



Process creation – Unix (example)

- ▶ `fork()` system call creates new process
 - Child process has a copy of the address space of the parent process
- ▶ `exec()` system call used after a `fork()`
 - To replace the process' memory space with a new program
- ▶ Parent process may call `wait()` to move itself off the ready queue until the child **terminates**

Forking Separate Process in C

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;


    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```

Operations on processes – Termination

- ▶ Process executes last statement, and then asks OS to delete it (`exit()`)
 - Returns status data from child to parent (via `wait()`)
 - Process' resources are deallocated by OS
- ▶ Parent may terminate the execution of children processes (`abort()`)
- ▶ Why?
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - If the parent terminates, the child also terminates (depends on OS)



Creating a Separate Process via Windows API

```
#include <stdio.h>
#include <windows.h>

int main(VOID)
{
    STARTUPINFO si;
    PROCESS_INFORMATION pi;


    /* allocate memory */
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));

    /* create child process */
    if (!CreateProcess(NULL, /* use command line */
        "C:\\WINDOWS\\system32\\mspaint.exe", /* command */
        NULL, /* don't inherit process handle */
        NULL, /* don't inherit thread handle */
        FALSE, /* disable handle inheritance */
        0, /* no creation flags */
        NULL, /* use parent's environment block */
        NULL, /* use parent's existing directory */
        &si,
        &pi))
    {
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
    /* parent will wait for the child to complete */
    WaitForSingleObject(pi.hProcess, INFINITE);
    printf("Child Complete");

    /* close handles */
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
}
```


Multiprocess architecture – Chrome (example)

- ▶ Many web browsers ran as single process (some still do)
- ▶ If one web site causes trouble, entire browser can hang or crash
- ▶ Can use multiprocess architecture to tackle this, e.g., **Chrome**
- ▶ 3 different types of processes
 - **Browser**: **only one**, created when Chrome starts
 - User interface, disk and network, etc.
 - **Renderer**: rendering webpages; generally one for each tab
 - Run in a **sandbox**
 - **Plug-in**: one for each type of plug-in (such as Flash or QuickTime) in use



Each tab represents a separate process.

Communications between processes

- ▶ Processes can be either **independent** or **cooperating**
- ▶ **Independent** processes
 - **Cannot** affect or be affected by the other processes
 - Does not share data with any other process
- ▶ **Cooperating** processes
 - **Can** affect or be affected by the other processes
 - Shares data with other processes
- ▶ Advantages of cooperation
 - Information sharing
 - Computation speed up
 - Modularity
 - Convenience

Communications between processes

- ▶ Cooperating processes need **interprocess communication** (IPC)
- ▶ Two models of IPC
 - **Shared memory**
 - A region of **memory** is established for sharing data
 - Information is exchanged by reading and writing to the shared memory
 - **Message passing**
 - Communication by exchanging **messages**

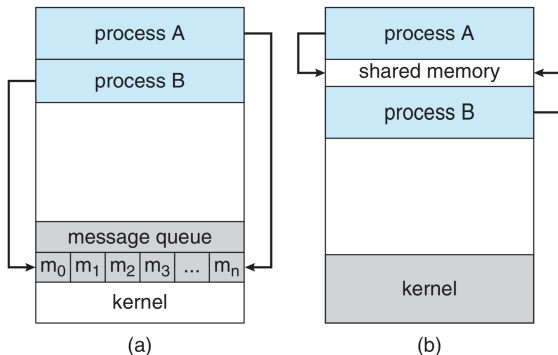


Figure 3.12 Communications models. (a) Message passing. (b) Shared memory.

Producer and consumer problem – Shared memory (1)

A common paradigm for cooperating processes

- ▶ A **producer** process produces information that is consumed by a **consumer** process

Shared-memory solution:

→ **Buffer**

- Bounded: **fixed** buffer size
 - Empty buffer: consumer must wait
 - Full buffer: producer must wait
- Unbounded: **no limit** on the buffer size
 - Empty buffer: consumer must wait
 - Producer can always produce

Example: bounded buffer

```
#define BUFFER_SIZE 10

typedef struct {
    ...
}item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

Producer and consumer problem – Shared memory (2)

One possible implementation

Producer

```
while (true) {  
    /* produce an item in next_produced */  
    while (((in + 1) % BUFFER_SIZE) == out)  
        ; /* do nothing */  
  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
}
```

- `next_produced` is a local variable

Consumer

```
item next_consumed;  
  
while (true) {  
    while (in == out)  
        ; /* do nothing */  
  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    /* consume the item in next_consumed */  
}
```

- This allows only at most `BUFFER_SIZE - 1` items in the buffer
- How do we have `BUFFER_SIZE` items?

- ▶ Provides a mechanism for process communicating and synchronising
- ▶ No shared memory
Useful in a distributed environment
- ▶ Provides at least two operations
 - `send(message)`
 - `receive(message)`
- ▶ Message size can be either **fixed** or **variable**

Message passing

- ▶ If processes P and Q want to communicate, they need to:
 - Establish a **communication link** between them, then
 - Exchange messages via send/receive operations
- ▶ We focus on **logical** implementation, not physical one (like bus or shared memory), possible methods:
 - Direct or indirect communication
 - Synchronous or asynchronous communication
 - Automatic or explicit buffering
- ▶ Implementation issues one needs to consider:
 - How are links established?
 - Can a link be associated with **more than two processes**?
 - How many links can there be between every pair of communicating processes?
 - What is the capacity of a link?
 - Is the size of a message that the link can accommodate fixed or variable?
 - Is a link unidirectional or bi-directional?

Message passing – Direct communication

- ▶ Processes must name each other **explicitly**
 - `send (P, message)` – send a message to P
 - `receive(Q, message)` – receive a message from Q
 - The above is called **symmetry addressing**
- ▶ For **asymmetry addressing**, only the sender names recipient
 - For recipient: use `receive(id, message)`
- ▶ Properties of communication link:
 - Link is established automatically
 - A link associate to **one and only one pair** of communicating processes
 - Exactly one link between each pair
 - The link can be unidirectional, but usually bi-directional

Message passing – Indirect communication (1)

- ▶ Messages are sent to and received from **mailboxes**, or **ports**.
- ▶ Each mailbox has a **unique identifier**
- ▶ Messages can be placed or removed from the mailbox by processes
- ▶ Two processes can communicate only if they have a shared mailbox
- ▶ `send()` and `receive()` are defined as
 - `send(A, message)` – send a message to mailbox A
 - `receive(A, message)` – receive a message from mailbox A
- ▶ Properties of communication link:
 - Link can only be established if processes share a mailbox
 - A link can associate to many processes
 - Each pair of processes can communicate via a number of different mailboxes
 - The link can be unidirectional, but usually bi-directional

Example: assume P_1 , P_2 and P_3 share the mailbox A

- ▶ P_1 sends a message to mailbox A
- ▶ P_2 and P_3 execute `receive()`
- ▶ Who will receive the message?
- ▶ Some possible answers:
 - Allow a link to associate with at most two processes
 - Allow at most one process to execute `receive()` at one time
 - Let the system choose **randomly** the recipient and notify the sender who it is

Message passing – Indirect communication (3)

A mailbox can be owned by either a **process** or **OS**

If a mailbox is owned by a **process**

- ▶ **Owner**: can only receive message through the mailbox
- ▶ **User**: can only send message to the mailbox
- ▶ If the owner process terminates, the mailbox will disappear

If a mailbox is owned by the **OS**

- ▶ The mailbox is not attached to any process
- ▶ OS provides a mechanism to processes to:
 - Create a new mailbox
 - Send and receive messages through a mailbox
 - Delete a mailbox
- ▶ Owner is the creating process by default
- ▶ Ownership can be passed to other processes \Rightarrow multiple receivers

Message passing can be either

- ▶ **Blocking** = Synchronous
 - **send**: sender is blocked until the message is received, either by receiving processes or mailbox
 - **receive**: receiver is blocked until a message is available
- ▶ **Nonblocking** = Asynchronous
 - **send**: sender resumes operation after sending a message
 - **receive**: receiver retrieves either a valid message or a null

Producer and consumer problem – Synchronisation

- ▶ When both `send()` and `receive()` are blocking
 - a **rendezvous** between the sender and the receiver
 - producer and consumer problem become **trivial**

Producer

```
message next_produced;  
while (true) {  
    /* produce an item in next_produced */  
  
    send(next_produced); /* send produced_element */  
}
```

Consumer

```
message next_consumed;  
while (true) {  
    receive(next_consumed); /* receive produced element */  
  
    /* consume the item in next_consumed */  
}
```


A message queue is required for messages

- ▶ **Zero capacity**: sender must wait for receiver (rendezvous)
- ▶ **Bounded capacity**: sender must wait if it is full
- ▶ **Unbounded capacity**: sender never waits

- ▶ Communication in ordinary producer-consumer fashion
 - Producer: writes to one end of the pipe (`write end`)
 - Receiver: reads from the other end (`read end`)
- ▶ **Unidirectional**
- ▶ **Two** pipes are needed for **bidirectional** communication

Ordinary pipe between parent and child in UNIX

- ▶ An ordinary pipe cannot be accessed from outside the process that created it
 - It is typical a **parent** process creates a pipe to communicate with its **child** process
- ▶ A pipe can be seen a special type of file
- ▶ In UNIX, a pipe that is accessed through the int **fd[]** file descriptors
- ▶ **fd[0]** is the read-end of the pipe, and **fd[1]** is the write-end




- ▶ More powerful than ordinary pipe
- ▶ Bidirectional
- ▶ Require no parent-child relationship
- ▶ One pipe can be used by several processes for communication

- ▶ Both shared memory and message passing can be used
- ▶ Also pipes
- ▶ Other strategies:
 - Sockets
 - Remote procedure calls (RPC)

Sockets

- ▶ Are defined as an **endpoint** for communication
 - ▶ Use client-server architecture
 - ▶ Two processes use a pair of sockets to communicate over a network
 - ▶ **Socket**: a concatenation of IP address and port
 - Port: A number included at start of message packet to differentiate network services on a host
- Socket example: 161.25.19.8 : 1625
- ▶ Ports under 1024 is **well-known** (or defined)



Read the socket example in Java in the textbook
(date server in Figure 3.27 & date client in Figure 3.28)

Remote procedure call (RPC) (1)

- ▶ Abstracts procedure calls between processes on networked systems
- ▶ Also uses **ports** to differentiate services
- ▶ Each message is addressed to an RPC daemon listening to a port on the remote system
- ▶ Each message contains
 - identifier specifying the function to execute, and
 - the parameters to pass to that function
- ▶ RPC system uses a **stub** on the client side to hide the details that enable the communications
- ▶ Typically, each procedure has a stub,
- ▶ When the client invokes a remote procedure
 - ① The RPC system calls the appropriate stub, and passes it the **parameters** for the remote procedure
 - ② The stub locates the **port** on the server and arranges the parameters

- ▶ Problem: Data representation can be different in the remote system
- ▶ Solution: interface definition language (IDL)
- ▶ Remote communication has higher potential to fail than local
 - Messages can be duplicated and executed more than once
 - One solution: OS has to ensure each message is executed **at most once**, or **exactly once**
- ▶ OS typically provides a **matchmaker** service to connect client and server

